

Project Summary:

Predict Housing Prices Project 1 Using Simple Linear Regression

Objective:

The task is to develop a simple linear regression model that predicts housing prices using the Boston Housing dataset. The entire process—data loading, preprocessing, model training, evaluation, and making predictions within a single script file.

Process:

1. Imported necessary libraries and connected to Kaggle API via Anaconda Prompt to download the dataset.
2. Extracted the data, defined file paths, and loaded the datasets into Jupyter Notebook.
3. Explored the data to understand its structure and contents.
4. Performed preprocessing:

```
Imputed missing numerical values with the mean.  
Imputed missing categorical values with the most frequent  
values using Simple Imputer.
```

5. Applied label encoding to categorical columns.
6. Detected outliers using the IQR method and visualized them.
7. Handled outliers by "Capping" (setting them to upper and lower bounds).
8. Conducted correlation analysis to assess relationships with the target variable.
9. Dropped constant columns and those with low correlation (threshold: 0.5).
10. Scaled the data using StandardScaler.
11. Split the data into training and test sets using `train_test_split`.
12. Trained a linear regression model on the training set and made predictions on the test set.
13. Evaluated model performance using MAE, RMSE, R^2 , and MSE.
14. Printed some sample predictions made by the model for comparison.

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error
from sklearn.preprocessing import StandardScaler
import seaborn as sns
import matplotlib.pyplot as plt
```

Downloaded the data using the Anaconda Prompt command

Unzipping the Data file

```
import zipfile
import os

# Define the path to the zip file and the directory where you want to
extract it
zip_file_path = r'C:\Users\ADDIS\house-prices-advanced-regression-
techniques.zip'
extract_dir = r'C:\Users\ADDIS\house-prices-data'

# Create the extraction directory
```

```
os.makedirs(extract_dir, exist_ok=True)

# Unzip the file
with zipfile.ZipFile(zip_file_path, 'r') as zip_ref:
    zip_ref.extractall(extract_dir)

print(f"Files extracted to {extract_dir}")

Files extracted to C:\Users\ADDIS\house-prices-data
```

Defining the path for the data

```
train_file_path = 'C:\\Users\\ADDIS\\house-prices-advanced-regression-techniques\\train.csv'
test_file_path = 'C:\\Users\\ADDIS\\house-prices-advanced-regression-techniques\\test.csv'
```

Loading the datasets to a dataframe

```
# Loading the datasets
train_data = pd.read_csv(train_file_path)
test_data = pd.read_csv(test_file_path)

# Display Data Types of Each Column
for column in train_data.columns:
    print(f"{column}: {train_data[column].dtype}")

Id: int64
MSSubClass: int64
MSZoning: object
LotFrontage: float64
LotArea: int64
Street: object
Alley: object
LotShape: object
LandContour: object
Utilities: object
LotConfig: object
LandSlope: object
Neighborhood: object
Condition1: object
Condition2: object
BldgType: object
HouseStyle: object
OverallQual: int64
OverallCond: int64
YearBuilt: int64
YearRemodAdd: int64
RoofStyle: object
```

RoofMatl: object
Exterior1st: object
Exterior2nd: object
MasVnrType: object
MasVnrArea: float64
ExterQual: object
ExterCond: object
Foundation: object
BsmtQual: object
BsmtCond: object
BsmtExposure: object
BsmtFinType1: object
BsmtFinSF1: int64
BsmtFinType2: object
BsmtFinSF2: int64
BsmtUnfSF: int64
TotalBsmtSF: int64
Heating: object
HeatingQC: object
CentralAir: object
Electrical: object
1stFlrSF: int64
2ndFlrSF: int64
LowQualFinSF: int64
GrLivArea: int64
BsmtFullBath: int64
BsmtHalfBath: int64
FullBath: int64
HalfBath: int64
BedroomAbvGr: int64
KitchenAbvGr: int64
KitchenQual: object
TotRmsAbvGrd: int64
Functional: object
Fireplaces: int64
FireplaceQu: object
GarageType: object
GarageYrBlt: float64
GarageFinish: object
GarageCars: int64
GarageArea: int64
GarageQual: object
GarageCond: object
PavedDrive: object
WoodDeckSF: int64
OpenPorchSF: int64
EnclosedPorch: int64
3SsnPorch: int64
ScreenPorch: int64

```
PoolArea: int64
PoolQC: object
Fence: object
MiscFeature: object
MiscVal: int64
MoSold: int64
YrSold: int64
SaleType: object
SaleCondition: object
SalePrice: int64
```

Summary Statistics for Numerical Columns in dataset

Check for Missing Values for train data

Missing Values in Each Column for test_data

```
print(train_data.shape)
print(test_data.shape)
```

```
(1460, 81)
(1459, 80)
```

```
train_data.head()
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape
0	1	60	RL	65.0	8450	Pave	NaN	Reg
1	2	20	RL	80.0	9600	Pave	NaN	Reg
2	3	60	RL	68.0	11250	Pave	NaN	IR1
3	4	70	RL	60.0	9550	Pave	NaN	IR1
4	5	60	RL	84.0	14260	Pave	NaN	IR1

	LandContour	Utilities	LotConfig	LandSlope	Neighborhood	Condition1
0	Lvl	AllPub	Inside	Gtl	CollgCr	Norm
1	Lvl	AllPub	FR2	Gtl	Veenker	Feedr
2	Lvl	AllPub	Inside	Gtl	CollgCr	Norm
3	Lvl	AllPub	Corner	Gtl	Crawfor	Norm
4	Lvl	AllPub	FR2	Gtl	NoRidge	Norm

	Condition2	BldgType	HouseStyle	OverallQual	OverallCond	YearBuilt
--	------------	----------	------------	-------------	-------------	-----------

0	Norm	1Fam	2Story	7	5	2003
1	Norm	1Fam	1Story	6	8	1976
2	Norm	1Fam	2Story	7	5	2001
3	Norm	1Fam	2Story	7	5	1915
4	Norm	1Fam	2Story	8	5	2000

	YearRemodAdd	RoofStyle	RoofMatl	Exterior1st	Exterior2nd	MasVnrType
0	2003	Gable	CompShg	VinylSd	VinylSd	BrkFace
1	1976	Gable	CompShg	MetalSd	MetalSd	NaN
2	2002	Gable	CompShg	VinylSd	VinylSd	BrkFace
3	1970	Gable	CompShg	Wd Sdng	Wd Shng	NaN
4	2000	Gable	CompShg	VinylSd	VinylSd	BrkFace

	MasVnrArea	ExterQual	ExterCond	Foundation	BsmtQual	BsmtCond
BsmtExposure \						
0	196.0	Gd	TA	PConc	Gd	TA
No						
1	0.0	TA	TA	CBlock	Gd	TA
Gd						
2	162.0	Gd	TA	PConc	Gd	TA
Mn						
3	0.0	TA	TA	BrkTil	TA	Gd
No						
4	350.0	Gd	TA	PConc	Gd	TA
Av						

BsmtFinType1	BsmtFinSF1	BsmtFinType2	BsmtFinSF2	BsmtUnfSF
0	706	Unf	0	150
1	978	Unf	0	284
2	486	Unf	0	434
3	216	Unf	0	540
4	655	Unf	0	490

Heating	HeatingQC	CentralAir	Electrical	1stFlrSF	2ndFlrSF
---------	-----------	------------	------------	----------	----------

LowQualFinSF	\					
0	GasA	Ex	Y	SBrkr	856	854
0						
1	GasA	Ex	Y	SBrkr	1262	0
0						
2	GasA	Ex	Y	SBrkr	920	866
0						
3	GasA	Gd	Y	SBrkr	961	756
0						
4	GasA	Ex	Y	SBrkr	1145	1053
0						
	GrLivArea	BsmtFullBath	BsmtHalfBath	FullBath	HalfBath	
BedroomAbvGr	\					
0	1710	1	0	2	1	
3						
1	1262	0	1	2	0	
3						
2	1786	1	0	2	1	
3						
3	1717	1	0	1	0	
3						
4	2198	1	0	2	1	
4						
	KitchenAbvGr	KitchenQual	TotRmsAbvGrd	Functional	Fireplaces	
FireplaceQu	\					
0	1	Gd	8	Typ	0	
NaN						
1	1	TA	6	Typ	1	
TA						
2	1	Gd	6	Typ	1	
TA						
3	1	Gd	7	Typ	1	
Gd						
4	1	Gd	9	Typ	1	
TA						
	GarageType	GarageYrBlt	GarageFinish	GarageCars	GarageArea	
GarageQual	\					
0	Attchd	2003.0	RFn	2	548	
TA						
1	Attchd	1976.0	RFn	2	460	
TA						
2	Attchd	2001.0	RFn	2	608	
TA						
3	Detchd	1998.0	Unf	3	642	
TA						
4	Attchd	2000.0	RFn	3	836	
TA						

	GarageCond	PavedDrive	WoodDeckSF	OpenPorchSF	EnclosedPorch
3SsnPorch \					
0	TA	Y	0	61	0
0					
1	TA	Y	298	0	0
0					
2	TA	Y	0	42	0
0					
3	TA	Y	0	35	272
0					
4	TA	Y	192	84	0
0					

	ScreenPorch	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	MoSold
YrSold \							
0	0	0	NaN	NaN	NaN	0	2
2008							
1	0	0	NaN	NaN	NaN	0	5
2007							
2	0	0	NaN	NaN	NaN	0	9
2008							
3	0	0	NaN	NaN	NaN	0	2
2006							
4	0	0	NaN	NaN	NaN	0	12
2008							

	SaleType	SaleCondition	SalePrice
0	WD	Normal	208500
1	WD	Normal	181500
2	WD	Normal	223500
3	WD	Abnorml	140000
4	WD	Normal	250000

Checks to see the columns and the type of data in each column (TEST DATA)

It showed that the data contains categorical and numerical columns which isnt good for prediction

test_data.head()

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley
LotShape \							
0	1461	20	RH	80.0	11622	Pave	NaN
Reg							
1	1462	20	RL	81.0	14267	Pave	NaN
IR1							
2	1463	60	RL	74.0	13830	Pave	NaN
IR1							
3	1464	60	RL	78.0	9978	Pave	NaN

IR1
 4 1465 120 RL 43.0 5005 Pave NaN
 IR1

	LandContour	Utilities	LotConfig	LandSlope	Neighborhood	Condition1	\
0	Lvl	AllPub	Inside	Gtl	NAmes	Feedr	
1	Lvl	AllPub	Corner	Gtl	NAmes	Norm	
2	Lvl	AllPub	Inside	Gtl	Gilbert	Norm	
3	Lvl	AllPub	Inside	Gtl	Gilbert	Norm	
4	HLS	AllPub	Inside	Gtl	StoneBr	Norm	

	Condition2	BldgType	HouseStyle	OverallQual	OverallCond	YearBuilt	\
0	Norm	1Fam	1Story	5	6	1961	
1	Norm	1Fam	1Story	6	6	1958	
2	Norm	1Fam	2Story	5	5	1997	
3	Norm	1Fam	2Story	6	6	1998	
4	Norm	TwnhsE	1Story	8	5	1992	

	YearRemodAdd	RoofStyle	RoofMatl	Exterior1st	Exterior2nd	MasVnrType	\
0	1961	Gable	CompShg	VinylSd	VinylSd	NaN	
1	1958	Hip	CompShg	Wd Sdng	Wd Sdng	BrkFace	
2	1998	Gable	CompShg	VinylSd	VinylSd	NaN	
3	1998	Gable	CompShg	VinylSd	VinylSd	BrkFace	
4	1992	Gable	CompShg	HdBoard	HdBoard	NaN	

	MasVnrArea	ExterQual	ExterCond	Foundation	BsmtQual	BsmtCond	BsmtExposure	\
0	0.0	TA	TA	CBlock	TA	TA	No	
1	108.0	TA	TA	CBlock	TA	TA	No	
2	0.0	TA	TA	PConc	Gd	TA	No	
3	20.0	TA	TA	PConc	TA	TA	No	
4	0.0	Gd	TA	PConc	Gd	TA	No	

BsmtFinType1	BsmtFinSF1	BsmtFinType2	BsmtFinSF2	BsmtUnfSF
--------------	------------	--------------	------------	-----------

TotalBsmtSF	\				
0	Rec	468.0	LwQ	144.0	270.0
882.0					
1	ALQ	923.0	Unf	0.0	406.0
1329.0					
2	GLQ	791.0	Unf	0.0	137.0
928.0					
3	GLQ	602.0	Unf	0.0	324.0
926.0					
4	ALQ	263.0	Unf	0.0	1017.0
1280.0					

	Heating	HeatingQC	CentralAir	Electrical	1stFlrSF	2ndFlrSF
LowQualFinSF	\					
0	GasA	TA	Y	SBrkr	896	0
0						
1	GasA	TA	Y	SBrkr	1329	0
0						
2	GasA	Gd	Y	SBrkr	928	701
0						
3	GasA	Ex	Y	SBrkr	926	678
0						
4	GasA	Ex	Y	SBrkr	1280	0
0						

GrLivArea	BsmtFullBath	BsmtHalfBath	FullBath	HalfBath
BedroomAbvGr \				
0 896	0.0	0.0	1	0
2				
1 1329	0.0	0.0	1	1
3				
2 1629	0.0	0.0	2	1
3				
3 1604	0.0	0.0	2	1
3				
4 1280	0.0	0.0	2	0
2				

KitchenAbvGr	KitchenQual	TotRmsAbvGrd	Functional	Fireplaces	
FireplaceQu \					
0 NaN	1	TA	5	Typ	0
1 NaN	1	Gd	6	Typ	0
2 TA	1	TA	6	Typ	1
3 Gd	1	Gd	7	Typ	1
4 NaN	1	Gd	5	Typ	0

	GarageType	GarageYrBlt	GarageFinish	GarageCars	GarageArea
0	Attchd	1961.0	Unf	1.0	730.0
TA					
1	Attchd	1958.0	Unf	1.0	312.0
TA					
2	Attchd	1997.0	Fin	2.0	482.0
TA					
3	Attchd	1998.0	Fin	2.0	470.0
TA					
4	Attchd	1992.0	RFn	2.0	506.0
TA					

	GarageCond	PavedDrive	WoodDeckSF	OpenPorchSF	EnclosedPorch
0	TA	Y	140	0	0
0					
1	TA	Y	393	36	0
0					
2	TA	Y	212	34	0
0					
3	TA	Y	360	36	0
0					
4	TA	Y	0	82	0
0					

	ScreenPorch	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	MoSold
0	120	0	NaN	MnPrv	NaN	0	6
2010							
1	0	0	NaN	NaN	Gar2	12500	6
2010							
2	0	0	NaN	MnPrv	NaN	0	3
2010							
3	0	0	NaN	NaN	NaN	0	6
2010							
4	144	0	NaN	NaN	NaN	0	1
2010							

	SaleType	SaleCondition
0	WD	Normal
1	WD	Normal
2	WD	Normal
3	WD	Normal
4	WD	Normal

Checks for missing values (TRAIN DATA)

Results showed that the test_data contains columns with missing values

```
train_data.isnull().sum()
```

```
Id                0
MSSubClass        0
MSZoning          0
LotFrontage      259
LotArea          0
...
MoSold           0
YrSold           0
SaleType         0
SaleCondition    0
SalePrice        0
Length: 81, dtype: int64
```

```
# Set display options to show all columns
```

```
pd.set_option('display.max_columns', None)
```

```
pd.set_option('display.max_rows', None)
```

```
# Display the count of missing values in each column
```

```
print(train_data.isnull().sum())
```

```
Id                0
MSSubClass        0
MSZoning          0
LotFrontage      259
LotArea          0
Street           0
Alley           1369
LotShape         0
LandContour      0
Utilities        0
LotConfig        0
LandSlope        0
Neighborhood     0
Condition1       0
Condition2       0
BldgType         0
HouseStyle       0
OverallQual      0
OverallCond      0
YearBuilt        0
YearRemodAdd     0
RoofStyle        0
RoofMatl         0
Exterior1st      0
Exterior2nd      0
```

MasVnrType	872
MasVnrArea	8
ExterQual	0
ExterCond	0
Foundation	0
BsmtQual	37
BsmtCond	37
BsmtExposure	38
BsmtFinType1	37
BsmtFinSF1	0
BsmtFinType2	38
BsmtFinSF2	0
BsmtUnfSF	0
TotalBsmtSF	0
Heating	0
HeatingQC	0
CentralAir	0
Electrical	1
1stFlrSF	0
2ndFlrSF	0
LowQualFinSF	0
GrLivArea	0
BsmtFullBath	0
BsmtHalfBath	0
FullBath	0
HalfBath	0
BedroomAbvGr	0
KitchenAbvGr	0
KitchenQual	0
TotRmsAbvGrd	0
Functional	0
Fireplaces	0
FireplaceQu	690
GarageType	81
GarageYrBlt	81
GarageFinish	81
GarageCars	0
GarageArea	0
GarageQual	81
GarageCond	81
PavedDrive	0
WoodDeckSF	0
OpenPorchSF	0
EnclosedPorch	0
3SsnPorch	0
ScreenPorch	0
PoolArea	0
PoolQC	1453
Fence	1179

```
MiscFeature      1406
MiscVal           0
MoSold            0
YrSold            0
SaleType          0
SaleCondition     0
SalePrice         0
dtype: int64
```

Checks for missing values (TEST DATA)

Results showed that the test_data contains columns with missing values

```
test_data.isnull().sum()
```

```
Id                0
MSSubClass        0
MSZoning          4
LotFrontage      227
LotArea           0
Street            0
Alley            1352
LotShape          0
LandContour       0
Utilities         2
LotConfig         0
LandSlope         0
Neighborhood      0
Condition1        0
Condition2        0
BldgType          0
HouseStyle        0
OverallQual       0
OverallCond       0
YearBuilt         0
YearRemodAdd      0
RoofStyle         0
RoofMatl          0
Exterior1st       1
Exterior2nd       1
MasVnrType        894
MasVnrArea        15
ExterQual         0
ExterCond         0
Foundation        0
BsmtQual          44
BsmtCond          45
BsmtExposure      44
```

BsmtFinType1	42
BsmtFinSF1	1
BsmtFinType2	42
BsmtFinSF2	1
BsmtUnfSF	1
TotalBsmtSF	1
Heating	0
HeatingQC	0
CentralAir	0
Electrical	0
1stFlrSF	0
2ndFlrSF	0
LowQualFinSF	0
GrLivArea	0
BsmtFullBath	2
BsmtHalfBath	2
FullBath	0
HalfBath	0
BedroomAbvGr	0
KitchenAbvGr	0
KitchenQual	1
TotRmsAbvGrd	0
Functional	2
Fireplaces	0
FireplaceQu	730
GarageType	76
GarageYrBlt	78
GarageFinish	78
GarageCars	1
GarageArea	1
GarageQual	78
GarageCond	78
PavedDrive	0
WoodDeckSF	0
OpenPorchSF	0
EnclosedPorch	0
3SsnPorch	0
ScreenPorch	0
PoolArea	0
PoolQC	1456
Fence	1169
MiscFeature	1408
MiscVal	0
MoSold	0
YrSold	0
SaleType	1
SaleCondition	0
dtype:	int64

```
# Set display options to show all columns
pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', None)

# Display the count of missing values in each column
print(test_data.isnull().sum())
```

Id	0
MSSubClass	0
MSZoning	4
LotFrontage	227
LotArea	0
Street	0
Alley	1352
LotShape	0
LandContour	0
Utilities	2
LotConfig	0
LandSlope	0
Neighborhood	0
Condition1	0
Condition2	0
BldgType	0
HouseStyle	0
OverallQual	0
OverallCond	0
YearBuilt	0
YearRemodAdd	0
RoofStyle	0
RoofMatl	0
Exterior1st	1
Exterior2nd	1
MasVnrType	894
MasVnrArea	15
ExterQual	0
ExterCond	0
Foundation	0
BsmtQual	44
BsmtCond	45
BsmtExposure	44
BsmtFinType1	42
BsmtFinSF1	1
BsmtFinType2	42
BsmtFinSF2	1
BsmtUnfSF	1
TotalBsmtSF	1
Heating	0
HeatingQC	0
CentralAir	0
Electrical	0

1stFlrSF	0
2ndFlrSF	0
LowQualFinSF	0
GrLivArea	0
BsmtFullBath	2
BsmtHalfBath	2
FullBath	0
HalfBath	0
BedroomAbvGr	0
KitchenAbvGr	0
KitchenQual	1
TotRmsAbvGrd	0
Functional	2
Fireplaces	0
FireplaceQu	730
GarageType	76
GarageYrBlt	78
GarageFinish	78
GarageCars	1
GarageArea	1
GarageQual	78
GarageCond	78
PavedDrive	0
WoodDeckSF	0
OpenPorchSF	0
EnclosedPorch	0
3SsnPorch	0
ScreenPorch	0
PoolArea	0
PoolQC	1456
Fence	1169
MiscFeature	1408
MiscVal	0
MoSold	0
YrSold	0
SaleType	1
SaleCondition	0

dtype: int64

Data Preprocessing

Data preprocessing is a crucial step in preparing a dataset for machine learning models. It ensures that the data is clean, consistent, and ready for analysis.

HANDLING MISSING VALUES IN NUMERICAL COLUMNS WITH THE MEAN OF THE COLUMN

```

from sklearn.impute import SimpleImputer

# Create the imputer object with mean strategy
imputer = SimpleImputer(strategy='mean')

# Select numerical columns excluding the target variable in train_data
numerical_cols =
train_data.drop(columns=['SalePrice']).select_dtypes(include=['float64', 'int64']).columns

# Fit the imputer on the training data excluding the target variable
imputer.fit(train_data[numerical_cols])

# Transform the training data excluding the target variable
train_data[numerical_cols] =
imputer.transform(train_data[numerical_cols])

# Transform the test data
test_data[numerical_cols] =
imputer.transform(test_data[numerical_cols])

```

HANDLING MISSING CATEGORICAL VALES IN THE DATASETS

```

# Fill missing categorical values with the nearest value using
backward fill
train_data = train_data.bfill()
test_data = test_data.bfill()

# Verify that there are no missing values in the categorical columns
print(train_data.isnull().sum())
print(test_data.isnull().sum())

```

Id	0
MSSubClass	0
MSZoning	0
LotFrontage	0
LotArea	0
Street	0
Alley	5
LotShape	0
LandContour	0
Utilities	0
LotConfig	0
LandSlope	0
Neighborhood	0
Condition1	0
Condition2	0
BldgType	0
HouseStyle	0

OverallQual	0
OverallCond	0
YearBuilt	0
YearRemodAdd	0
RoofStyle	0
RoofMatl	0
Exterior1st	0
Exterior2nd	0
MasVnrType	3
MasVnrArea	0
ExterQual	0
ExterCond	0
Foundation	0
BsmtQual	0
BsmtCond	0
BsmtExposure	0
BsmtFinType1	0
BsmtFinSF1	0
BsmtFinType2	0
BsmtFinSF2	0
BsmtUnfSF	0
TotalBsmtSF	0
Heating	0
HeatingQC	0
CentralAir	0
Electrical	0
1stFlrSF	0
2ndFlrSF	0
LowQualFinSF	0
GrLivArea	0
BsmtFullBath	0
BsmtHalfBath	0
FullBath	0
HalfBath	0
BedroomAbvGr	0
KitchenAbvGr	0
KitchenQual	0
TotRmsAbvGrd	0
Functional	0
Fireplaces	0
FireplaceQu	2
GarageType	0
GarageYrBlt	0
GarageFinish	0
GarageCars	0
GarageArea	0
GarageQual	0
GarageCond	0
PavedDrive	0

WoodDeckSF	0
OpenPorchSF	0
EnclosedPorch	0
3SsnPorch	0
ScreenPorch	0
PoolArea	0
PoolQC	36
Fence	2
MiscFeature	2
MiscVal	0
MoSold	0
YrSold	0
SaleType	0
SaleCondition	0
SalePrice	0
dtype: int64	
Id	0
MSSubClass	0
MSZoning	0
LotFrontage	0
LotArea	0
Street	0
Alley	47
LotShape	0
LandContour	0
Utilities	0
LotConfig	0
LandSlope	0
Neighborhood	0
Condition1	0
Condition2	0
BldgType	0
HouseStyle	0
OverallQual	0
OverallCond	0
YearBuilt	0
YearRemodAdd	0
RoofStyle	0
RoofMatl	0
Exterior1st	0
Exterior2nd	0
MasVnrType	0
MasVnrArea	0
ExterQual	0
ExterCond	0
Foundation	0
BsmtQual	0
BsmtCond	0
BsmtExposure	0

BsmtFinType1	0
BsmtFinSF1	0
BsmtFinType2	0
BsmtFinSF2	0
BsmtUnfSF	0
TotalBsmtSF	0
Heating	0
HeatingQC	0
CentralAir	0
Electrical	0
1stFlrSF	0
2ndFlrSF	0
LowQualFinSF	0
GrLivArea	0
BsmtFullBath	0
BsmtHalfBath	0
FullBath	0
HalfBath	0
BedroomAbvGr	0
KitchenAbvGr	0
KitchenQual	0
TotRmsAbvGrd	0
Functional	0
Fireplaces	0
FireplaceQu	0
GarageType	0
GarageYrBlt	0
GarageFinish	0
GarageCars	0
GarageArea	0
GarageQual	0
GarageCond	0
PavedDrive	0
WoodDeckSF	0
OpenPorchSF	0
EnclosedPorch	0
3SsnPorch	0
ScreenPorch	0
PoolArea	0
PoolQC	208
Fence	1
MiscFeature	1
MiscVal	0
MoSold	0
YrSold	0
SaleType	0
SaleCondition	0
dtype:	int64

CONFIRMATION THAT ALL MISSING NUMERICAL VALUES ARE HANDLED

The code handled the missing numerical values appropriately

```
import pandas as pd

# Select numerical columns in train_data excluding the target variable
numerical_cols_train =
train_data.drop(columns=['SalePrice']).select_dtypes(include=['float64',
'int64']).columns

# Check for null values in numerical columns of train_data
null_columns_train = train_data[numerical_cols_train].isnull().sum()
null_columns_train = null_columns_train[null_columns_train > 0]
print("Numerical columns with null values in train_data:")
print(null_columns_train)

# Select numerical columns in test_data
numerical_cols_test = test_data.select_dtypes(include=['float64',
'int64']).columns

# Check for null values in numerical columns of test_data
null_columns_test = test_data[numerical_cols_test].isnull().sum()
null_columns_test = null_columns_test[null_columns_test > 0]
print("Numerical columns with null values in test_data:")
print(null_columns_test)

Numerical columns with null values in train_data:
Series([], dtype: int64)
Numerical columns with null values in test_data:
Series([], dtype: int64)
```

CONFIRMATION OF THE BACKFILLING OF CATEGORICAL COLUMNS

#It happend that the backfill of categorical columns did not fill the null values because the missing values are scattered randomly rather than in a sequence, backfilling might not fill all gaps correctly.

```
import pandas as pd

# Select categorical columns in train_data excluding the target variable
categorical_cols_train =
train_data.drop(columns=['SalePrice']).select_dtypes(include=['object']
).columns

# Check for null values in categorical columns of train_data
null_columns_train = train_data[categorical_cols_train].isnull().sum()
null_columns_train = null_columns_train[null_columns_train > 0]
```

```

print("Categorical columns with null values in train_data:")
print(null_columns_train)

# Select categorical columns in test_data
categorical_cols_test =
test_data.select_dtypes(include=['object']).columns

# Check for null values in categorical columns of test_data
null_columns_test = test_data[categorical_cols_test].isnull().sum()
null_columns_test = null_columns_test[null_columns_test > 0]
print("Categorical columns with null values in test_data:")
print(null_columns_test)

Categorical columns with null values in train_data:
Alley          5
MasVnrType     3
FireplaceQu    2
PoolQC        36
Fence          2
MiscFeature    2
dtype: int64
Categorical columns with null values in test_data:
Alley          47
PoolQC        208
Fence          1
MiscFeature    1
dtype: int64

## CODE to fill the categorical columns

from sklearn.impute import SimpleImputer

# Create the imputer object with the most frequent strategy
imputer = SimpleImputer(strategy='most_frequent')

# Select categorical columns excluding the target variable in
train_data
categorical_cols_train =
train_data.drop(columns=['SalePrice']).select_dtypes(include=['object'
]).columns

# Fit the imputer on the training data excluding the target variable
imputer.fit(train_data[categorical_cols_train])

# Transform the training data excluding the target variable
train_data[categorical_cols_train] =
imputer.transform(train_data[categorical_cols_train])

# Select categorical columns in test_data

```

```

categorical_cols_test =
test_data.select_dtypes(include=['object']).columns

# Transform the test data
test_data[categorical_cols_test] =
imputer.transform(test_data[categorical_cols_test])

## confirmation of the filling of missing categorical columns

import pandas as pd

# Select categorical columns in train_data excluding the target
variable
categorical_cols_train =
train_data.drop(columns=['SalePrice']).select_dtypes(include=['object'
]).columns

# Check for null values in categorical columns of train_data
null_columns_train = train_data[categorical_cols_train].isnull().sum()
null_columns_train = null_columns_train[null_columns_train > 0]
print("Categorical columns with null values in train_data:")
print(null_columns_train)

# Select categorical columns in test_data
categorical_cols_test =
test_data.select_dtypes(include=['object']).columns

# Check for null values in categorical columns of test_data
null_columns_test = test_data[categorical_cols_test].isnull().sum()
null_columns_test = null_columns_test[null_columns_test > 0]
print("Categorical columns with null values in test_data:")
print(null_columns_test)

Categorical columns with null values in train_data:
Series([], dtype: int64)
Categorical columns with null values in test_data:
Series([], dtype: int64)

```

Label encoding categorical columns to Numerical data so it can fit into the model

```

from sklearn.preprocessing import LabelEncoder

# Initialize the LabelEncoder
label_encoder = LabelEncoder()

# Combine train and test data for fitting the encoder
combined_data = pd.concat([train_data.drop(columns=['SalePrice']),

```



```
test_data])

# Select categorical columns in the combined data
categorical_cols =
combined_data.select_dtypes(include=['object']).columns

# Apply label encoding to each categorical column in the combined data
for col in categorical_cols:
    label_encoder.fit(combined_data[col].astype(str))
    train_data[col] =
label_encoder.transform(train_data[col].astype(str))
    test_data[col] =
label_encoder.transform(test_data[col].astype(str))

# Verify the changes
print("First few rows of the transformed train_data:")
print(train_data.head())

print("First few rows of the transformed test_data:")
print(test_data.head())
```

First few rows of the transformed train_data:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley
0	1.0	60.0	3	65.0	8450.0	1	0
1	2.0	20.0	3	80.0	9600.0	1	0
2	3.0	60.0	3	68.0	11250.0	1	0
3	4.0	70.0	3	60.0	9550.0	1	0
4	5.0	60.0	3	84.0	14260.0	1	0

	LandContour	Utilities	LotConfig	LandSlope	Neighborhood
0	3	0	4	0	20
1	3	0	2	0	17
2	3	0	4	0	20
3	3	0	0	0	21
4	3	0	2	0	7

	Condition2	BldgType	HouseStyle	OverallQual	OverallCond
0	3	0	4	0	20
1	3	0	2	0	17
2	3	0	4	0	20
3	3	0	0	0	21
4	3	0	2	0	7

0	2	0	10	7.0	5.0
2003.0					
1	2	0	5	6.0	8.0
1976.0					
2	2	0	10	7.0	5.0
2001.0					
3	2	0	10	7.0	5.0
1915.0					
4	2	0	10	8.0	5.0
2000.0					

	YearRemodAdd	RoofStyle	RoofMatl	Exterior1st	Exterior2nd
MasVnrType \					
0	2003.0	1	1	4	5
1					
1	1976.0	1	1	13	14
1					
2	2002.0	1	1	4	5
1					
3	1970.0	1	1	5	7
1					
4	2000.0	1	1	4	5
1					

	MasVnrArea	ExterQual	ExterCond	Foundation	BsmtQual	BsmtCond	\
0	196.0	2	4	2	2	3	
1	0.0	3	4	1	2	3	
2	162.0	2	4	2	2	3	
3	0.0	3	4	0	3	1	
4	350.0	2	4	2	2	3	

	BsmtExposure	BsmtFinType1	BsmtFinSF1	BsmtFinType2	BsmtFinSF2	\
0	3	2	706.0	5	0.0	
1	1	0	978.0	5	0.0	
2	2	2	486.0	5	0.0	
3	3	0	216.0	5	0.0	
4	0	2	655.0	5	0.0	

	BsmtUnfSF	TotalBsmtSF	Heating	HeatingQC	CentralAir	Electrical	\
0	150.0	856.0	1	0	1	4	
1	284.0	1262.0	1	0	1	4	
2	434.0	920.0	1	0	1	4	
3	540.0	756.0	1	2	1	4	
4	490.0	1145.0	1	0	1	4	

	1stFlrSF	2ndFlrSF	LowQualFinSF	GrLivArea	BsmtFullBath
BsmtHalfBath \					
0	856.0	854.0	0.0	1710.0	1.0
0.0					
1	1262.0	0.0	0.0	1262.0	0.0
1.0					
2	920.0	866.0	0.0	1786.0	1.0
0.0					
3	961.0	756.0	0.0	1717.0	1.0
0.0					
4	1145.0	1053.0	0.0	2198.0	1.0
0.0					

	FullBath	HalfBath	BedroomAbvGr	KitchenAbvGr	KitchenQual
TotRmsAbvGrd \					
0	2.0	1.0	3.0	1.0	2
8.0					
1	2.0	0.0	3.0	1.0	3
6.0					
2	2.0	1.0	3.0	1.0	2
6.0					
3	1.0	0.0	3.0	1.0	2
7.0					
4	2.0	1.0	4.0	1.0	2
9.0					

	Functional	Fireplaces	FireplaceQu	GarageType	GarageYrBlt
GarageFinish \					
0	6	0.0	4	1	2003.0
1					
1	6	1.0	4	1	1976.0
1					
2	6	1.0	4	1	2001.0
1					
3	6	1.0	2	6	1998.0
2					
4	6	1.0	4	1	2000.0
1					

	GarageCars	GarageArea	GarageQual	GarageCond	PavedDrive
WoodDeckSF \					
0	2.0	548.0	4	4	2
0.0					
1	2.0	460.0	4	4	2
298.0					
2	2.0	608.0	4	4	2
0.0					
3	3.0	642.0	4	4	2
0.0					

4	3.0	836.0	4	4	2
192.0					

	OpenPorchSF	EnclosedPorch	3SsnPorch	ScreenPorch	PoolArea
PoolQC \					
0	61.0	0.0	0.0	0.0	0.0
0					
1	0.0	0.0	0.0	0.0	0.0
0					
2	42.0	0.0	0.0	0.0	0.0
0					
3	35.0	272.0	0.0	0.0	0.0
0					
4	84.0	0.0	0.0	0.0	0.0
0					

	Fence	MiscFeature	MiscVal	MoSold	YrSold	SaleType
SaleCondition \						
0	2	2	0.0	2.0	2008.0	8
4						
1	2	2	0.0	5.0	2007.0	8
4						
2	2	2	0.0	9.0	2008.0	8
4						
3	2	2	0.0	2.0	2006.0	8
0						
4	2	2	0.0	12.0	2008.0	8
4						

	SalePrice
0	208500
1	181500
2	223500
3	140000
4	250000

First few rows of the transformed test_data:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street
Alley \						
0	1461.0	20.0	7	80.0	11622.0	3 3
1	1462.0	20.0	8	81.0	14267.0	3 3
2	1463.0	60.0	8	74.0	13830.0	3 3
3	1464.0	60.0	8	78.0	9978.0	3 3
4	1465.0	120.0	8	43.0	5005.0	3 3

LotShape	LandContour	Utilities	LotConfig	LandSlope
----------	-------------	-----------	-----------	-----------

Neighborhood \					
0	7	7	2	9	3
37					
1	4	7	2	5	3
37					
2	4	7	2	9	3
33					
3	4	7	2	9	3
33					
4	4	5	2	9	3
47					
Condition1	Condition2	BldgType	HouseStyle	OverallQual	
OverallCond \					
0	10	10	2	4	5.0
6.0					
1	11	10	2	4	6.0
6.0					
2	11	10	2	7	5.0
5.0					
3	11	10	2	7	6.0
6.0					
4	11	10	9	4	8.0
5.0					
YearBuilt	YearRemodAdd	RoofStyle	RoofMatl	Exterior1st	
Exterior2nd \					
0	1961.0	1961.0	7	8	25
28					
1	1958.0	1958.0	9	8	26
29					
2	1997.0	1998.0	7	8	25
28					
3	1998.0	1998.0	7	8	25
28					
4	1992.0	1992.0	7	8	21
22					
MasVnrType	MasVnrArea	ExterQual	ExterCond	Foundation	BsmtQual
\					
0	4	0.0	7	9	7
1	4	108.0	7	9	7
2	4	0.0	7	9	8
					6
3	4	20.0	7	9	8
					7
4	4	0.0	6	9	8
					6

BsmtCond	BsmtExposure	BsmtFinType1	BsmtFinSF1	BsmtFinType2	BsmtFinSF2
0	7	10	468.0	9	144.0
1	7	6	923.0	11	0.0
2	7	8	791.0	11	0.0
3	7	8	602.0	11	0.0
4	7	6	263.0	11	0.0

\	BsmtUnfSF	TotalBsmtSF	Heating	HeatingQC	CentralAir	Electrical
0	270.0	882.0	6	9	3	8
1	406.0	1329.0	6	9	3	8
2	137.0	928.0	6	7	3	8
3	324.0	926.0	6	5	3	8
4	1017.0	1280.0	6	5	3	8

1stFlrSF	2ndFlrSF	LowQualFinSF	GrLivArea	BsmtFullBath	BsmtHalfBath
0	896.0	0.0	0.0	896.0	0.0
1	1329.0	0.0	0.0	1329.0	0.0
2	928.0	701.0	0.0	1629.0	0.0
3	926.0	678.0	0.0	1604.0	0.0
4	1280.0	0.0	0.0	1280.0	0.0

FullBath	HalfBath	BedroomAbvGr	KitchenAbvGr	KitchenQual	TotRmsAbvGrd
0	1.0	0.0	2.0	1.0	7
1	1.0	1.0	3.0	1.0	6
2	2.0	1.0	3.0	1.0	7
3	2.0	1.0	3.0	1.0	6

4	2.0	0.0	2.0	1.0	6
5.0					

Functional GarageFinish	Fireplaces \	FireplaceQu	GarageType	GarageYrBl	GarageYrBl
0	13	0.0	9	7	1961.0
5					
1	13	0.0	9	7	1958.0
5					
2	13	1.0	9	7	1997.0
3					
3	13	1.0	7	7	1998.0
3					
4	13	0.0	9	7	1992.0
4					

GarageCars	GarageArea	GarageQual	GarageCond	PavedDrive
WoodDeckSF	\			
0	1.0	730.0	8	9
140.0				5
1	1.0	312.0	8	9
393.0				5
2	2.0	482.0	8	9
212.0				5
3	2.0	470.0	8	9
360.0				5
4	2.0	506.0	8	9
0.0				5

OpenPorchSF	EnclosedPorch	3SsnPorch	ScreenPorch	PoolArea
PoolQC	\			
0	0.0	0.0	0.0	120.0
3				0.0
1	36.0	0.0	0.0	0.0
3				0.0
2	34.0	0.0	0.0	0.0
3				0.0
3	36.0	0.0	0.0	0.0
3				0.0
4	82.0	0.0	0.0	144.0
3				0.0

Fence	MiscFeature	MiscVal	MoSold	YrSold	SaleType
SaleCondition					
0	6	4	0.0	6.0	2010.0
10					17
1	6	4	12500.0	6.0	2010.0
10					17
2	6	6	0.0	3.0	2010.0
10					17

3	4	6	0.0	6.0	2010.0	17
10						
4	4	6	0.0	1.0	2010.0	17
10						

OUTLIER DETECTION

it was detected that the both datasets contains outleirs which if not carefully handled, can result to skewed predictions

```
# Calculate Q1 (25th percentile) and Q3 (75th percentile)
```

```
Q1 = train_data.quantile(0.25)
```

```
Q3 = train_data.quantile(0.75)
```

```
IQR = Q3 - Q1
```

```
# Define outlier boundaries
```

```
lower_bound = Q1 - 1.5 * IQR
```

```
upper_bound = Q3 + 1.5 * IQR
```

```
# Identify outliers
```

```
outliers = ((train_data < lower_bound) | (train_data >
```

```
upper_bound)).any(axis=1)
```

```
outlier_indices = train_data[outliers].index
```

```
train_data.head()
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley
LotShape \							
0	1.0	60.0	3	65.0	8450.0	1	0
3							
1	2.0	20.0	3	80.0	9600.0	1	0
3							
2	3.0	60.0	3	68.0	11250.0	1	0
0							
3	4.0	70.0	3	60.0	9550.0	1	0
0							
4	5.0	60.0	3	84.0	14260.0	1	0
0							

	LandContour	Utilities	LotConfig	LandSlope	Neighborhood
Condition1 \					
0	3	0	4	0	20
2					

1	3	0	2	0	17
1					
2	3	0	4	0	20
2					
3	3	0	0	0	21
2					
4	3	0	2	0	7
2					

Condition2	BldgType	HouseStyle	OverallQual	OverallCond
YearBuilt \				
0	2	0	10	7.0
2003.0				5.0
1	2	0	5	6.0
1976.0				8.0
2	2	0	10	7.0
2001.0				5.0
3	2	0	10	7.0
1915.0				5.0
4	2	0	10	8.0
2000.0				5.0

YearRemodAdd	RoofStyle	RoofMatl	Exterior1st	Exterior2nd
MasVnrType \				
0	2003.0	1	1	4
1				5
1	1976.0	1	1	13
1				14
2	2002.0	1	1	4
1				5
3	1970.0	1	1	5
1				7
4	2000.0	1	1	4
1				5

MasVnrArea	ExterQual	ExterCond	Foundation	BsmtQual	BsmtCond	\
0	196.0	2	4	2	2	3
1	0.0	3	4	1	2	3
2	162.0	2	4	2	2	3
3	0.0	3	4	0	3	1
4	350.0	2	4	2	2	3

BsmtExposure	BsmtFinType1	BsmtFinSF1	BsmtFinType2	BsmtFinSF2	\
0	3	2	706.0	5	0.0
1	1	0	978.0	5	0.0
2	2	2	486.0	5	0.0
3	3	0	216.0	5	0.0
4	0	2	655.0	5	0.0

BsmtUnfSF	TotalBsmtSF	Heating	HeatingQC	CentralAir	Electrical
-----------	-------------	---------	-----------	------------	------------

\						
0	150.0	856.0	1	0	1	4
1	284.0	1262.0	1	0	1	4
2	434.0	920.0	1	0	1	4
3	540.0	756.0	1	2	1	4
4	490.0	1145.0	1	0	1	4
	1stFlrSF	2ndFlrSF	LowQualFinSF	GrLivArea	BsmtFullBath	
BsmtHalfBath	\					
0	856.0	854.0	0.0	1710.0	1.0	
0.0						
1	1262.0	0.0	0.0	1262.0	0.0	
1.0						
2	920.0	866.0	0.0	1786.0	1.0	
0.0						
3	961.0	756.0	0.0	1717.0	1.0	
0.0						
4	1145.0	1053.0	0.0	2198.0	1.0	
0.0						
	FullBath	HalfBath	BedroomAbvGr	KitchenAbvGr	KitchenQual	
TotRmsAbvGrd	\					
0	2.0	1.0	3.0	1.0	2	
8.0						
1	2.0	0.0	3.0	1.0	3	
6.0						
2	2.0	1.0	3.0	1.0	2	
6.0						
3	1.0	0.0	3.0	1.0	2	
7.0						
4	2.0	1.0	4.0	1.0	2	
9.0						
	Functional	Fireplaces	FireplaceQu	GarageType	GarageYrBlt	
GarageFinish	\					
0	6	0.0	4	1	2003.0	
1						
1	6	1.0	4	1	1976.0	
1						
2	6	1.0	4	1	2001.0	
1						
3	6	1.0	2	6	1998.0	
2						
4	6	1.0	4	1	2000.0	
1						

	GarageCars	GarageArea	GarageQual	GarageCond	PavedDrive
0	2.0	548.0	4	4	2
1	2.0	460.0	4	4	2
2	2.0	608.0	4	4	2
3	3.0	642.0	4	4	2
4	3.0	836.0	4	4	2

	OpenPorchSF	EnclosedPorch	3SsnPorch	ScreenPorch	PoolArea
0	61.0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0
2	42.0	0.0	0.0	0.0	0.0
3	35.0	272.0	0.0	0.0	0.0
4	84.0	0.0	0.0	0.0	0.0

	Fence	MiscFeature	MiscVal	MoSold	YrSold	SaleType
0	2	2	0.0	2.0	2008.0	8
1	2	2	0.0	5.0	2007.0	8
2	2	2	0.0	9.0	2008.0	8
3	2	2	0.0	2.0	2006.0	8
4	2	2	0.0	12.0	2008.0	8

	SalePrice
0	208500
1	181500
2	223500
3	140000
4	250000

outlier detection for train_dataset

to see columns with outliers in the train_data set

```
import pandas as pd

# Function to detect outliers using the IQR method
def detect_outliers(df):
    outliers = {}
    for column in df.select_dtypes(include=['int64',
'float64']).columns:
        Q1 = df[column].quantile(0.25)
        Q3 = df[column].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        outlier_indices = df[(df[column] < lower_bound) | (df[column]
> upper_bound)].index
        if not outlier_indices.empty:
            outliers[column] = len(outlier_indices)
    return outliers

# Detect outliers in the training data
outliers = detect_outliers(train_data)

# Print outliers in the desired format
print(f"{'column name':<25} {'outlier found':<15}")
for column, count in outliers.items():
    print(f"{'column':<25} {'count':<15}")
```

column name	outlier found
MSSubClass	103
LotFrontage	106
LotArea	69
OverallQual	2
OverallCond	125
YearBuilt	7
MasVnrArea	98
BsmtFinSF1	7
BsmtFinSF2	167
BsmtUnfSF	29
TotalBsmtSF	61
1stFlrSF	20
2ndFlrSF	2
LowQualFinSF	26
GrLivArea	31
BsmtFullBath	1
BsmtHalfBath	82
BedroomAbvGr	35
KitchenAbvGr	68
TotRmsAbvGrd	30

Fireplaces	5
GarageYrBltd	1
GarageCars	5
GarageArea	21
WoodDeckSF	32
OpenPorchSF	77
EnclosedPorch	208
3SsnPorch	24
ScreenPorch	116
PoolArea	7
MiscVal	52
SalePrice	61

outleir detection for test_data set

to see columns in the test_data with outliers

```
import pandas as pd
```

Function to detect outliers using the IQR method

```
def detect_outliers(df):
    outliers = {}
    for column in df.select_dtypes(include=['int64',
'float64']).columns:
        Q1 = df[column].quantile(0.25)
        Q3 = df[column].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        outlier_indices = df[(df[column] < lower_bound) | (df[column]
> upper_bound)].index
        if not outlier_indices.empty:
            outliers[column] = len(outlier_indices)
    return outliers
```

Detect outliers in the training data

```
outliers = detect_outliers(test_data)
```

Print outliers in the desired format

```
print(f"{'column name':<25} {'outlier found':<15}")
for column, count in outliers.items():
    print(f"{'column':<25} {'count':<15}")
```

column name	outlier found
MSSubClass	104
LotFrontage	141
LotArea	60
OverallQual	2
OverallCond	127

YearBuilt	2
MasVnrArea	104
BsmtFinSF1	8
BsmtFinSF2	181
BsmtUnfSF	26
TotalBsmtSF	61
1stFlrSF	23
2ndFlrSF	5
LowQualFinSF	14
GrLivArea	44
BsmtFullBath	1
BsmtHalfBath	95
FullBath	4
BedroomAbvGr	43
KitchenAbvGr	66
TotRmsAbvGrd	21
Fireplaces	7
GarageYrBlt	3
GarageCars	12
GarageArea	21
WoodDeckSF	35
OpenPorchSF	79
EnclosedPorch	251
3SsnPorch	13
ScreenPorch	140
PoolArea	6
MiscVal	51

HANDLING OUTLIERS (CAPPING)

“Capping” outliers means replacing extreme values with a specified limit to reduce their impact on the dataset. Instead of removing outliers, which can result in data loss, capping modifies the outliers to fall within a defined range. This approach retains all data points but adjusts the extreme values to be less influential.

Based on the lists provided, the FullBath column has outliers in the test_data but not in the train_data. The capping of the outliers will be handled separately with the columns in the lists .

```
def cap_outliers(df, columns):
    for column in columns:
        Q1 = df[column].quantile(0.25)
        Q3 = df[column].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        df[column] = df[column].apply(lambda x: lower_bound if x <
```

```

lower_bound else upper_bound if x > upper_bound else x)
    return df

# List of columns to cap in train_data
columns_to_cap = [
    'MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual',
    'OverallCond', 'YearBuilt',
    'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF',
    'TotalBsmtSF', '1stFlrSF',
    '2ndFlrSF', 'LowQualFinSF', 'GrLivArea', 'BsmtFullBath',
    'BsmtHalfBath', 'BedroomAbvGr',
    'KitchenAbvGr', 'TotRmsAbvGrd', 'Fireplaces', 'GarageYrBlt',
    'GarageCars', 'GarageArea',
    'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SsnPorch',
    'ScreenPorch', 'PoolArea',
    'MiscVal'
]

# Cap outliers in the specified columns of train_data
train_data_capped =
cap_outliers(train_data.drop(columns=['SalePrice']), columns_to_cap)

# Add the SalePrice column back to the capped training data
train_data_capped['SalePrice'] = train_data['SalePrice']

# Verify the changes
print(f"Original train_data shape: {train_data.shape}")
print(f"Capped train_data shape: {train_data_capped.shape}")

Original train_data shape: (1460, 81)
Capped train_data shape: (1460, 81)

def cap_outliers(df, columns):
    for column in columns:
        Q1 = df[column].quantile(0.25)
        Q3 = df[column].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        df[column] = df[column].apply(lambda x: lower_bound if x <
lower_bound else upper_bound if x > upper_bound else x)
    return df

# List of columns to cap in test_data
columns_to_cap_test = [
    'MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual',
    'OverallCond', 'YearBuilt',
    'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF',
    'TotalBsmtSF', '1stFlrSF',
    '2ndFlrSF', 'LowQualFinSF', 'GrLivArea', 'BsmtFullBath',

```

```

'BsmthalfBath', 'FullBath',
    'BedroomAbvGr', 'KitchenAbvGr', 'TotRmsAbvGrd', 'Fireplaces',
'GarageYrBlt', 'GarageCars',
    'GarageArea', 'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch',
'3SsnPorch', 'ScreenPorch',
    'PoolArea', 'MiscVal'
]

# Cap outliers in the specified columns of test_data
test_data_capped = cap_outliers(test_data, columns_to_cap_test)

# Verify the changes
print(f"Original test_data shape: {test_data.shape}")
print(f"Capped test_data shape: {test_data_capped.shape}")

Original test_data shape: (1459, 80)
Capped test_data shape: (1459, 80)

# Compare original and capped values for a few columns
columns_to_check = ['LotFrontage', 'GrLivArea', 'FullBath'] # Add
more columns as needed

for column in columns_to_check:
    original_values = test_data[column].head(10) # Display first 10
values for comparison
    capped_values = test_data_capped[column].head(10)
    print(f"Original {column} values:\n{original_values}")
    print(f"Capped {column} values:\n{capped_values}\n")

Original LotFrontage values:
0    80.000000
1    81.000000
2    74.000000
3    78.000000
4    43.000000
5    75.000000
6    70.049958
7    63.000000
8    85.000000
9    70.000000
Name: LotFrontage, dtype: float64
Capped LotFrontage values:
0    80.000000
1    81.000000
2    74.000000
3    78.000000
4    43.000000
5    75.000000
6    70.049958
7    63.000000

```



```
8      85.000000
9      70.000000
Name: LotFrontage, dtype: float64
```

Original GrLivArea values:

```
0      896.0
1     1329.0
2     1629.0
3     1604.0
4     1280.0
5     1655.0
6     1187.0
7     1465.0
8     1341.0
9      882.0
```

Name: GrLivArea, dtype: float64

Capped GrLivArea values:

```
0      896.0
1     1329.0
2     1629.0
3     1604.0
4     1280.0
5     1655.0
6     1187.0
7     1465.0
8     1341.0
9      882.0
```

Name: GrLivArea, dtype: float64

Original FullBath values:

```
0      1.0
1      1.0
2      2.0
3      2.0
4      2.0
5      2.0
6      2.0
7      2.0
8      1.0
9      1.0
```

Name: FullBath, dtype: float64

Capped FullBath values:

```
0      1.0
1      1.0
2      2.0
3      2.0
4      2.0
5      2.0
6      2.0
```

```
7     2.0
8     1.0
9     1.0
```

```
Name: FullBath, dtype: float64
```

```
# Summary statistics before capping
print("Summary statistics before capping:")
print(test_data[columns_to_check].describe())
```

```
# Summary statistics after capping
print("Summary statistics after capping:")
print(test_data_capped[columns_to_check].describe())
```

Summary statistics before capping:

	LotFrontage	GrLivArea	FullBath
count	1459.000000	1459.000000	1459.000000
mean	68.423126	1478.000685	1.569568
std	17.164354	457.873870	0.549778
min	33.000000	407.000000	0.000000
25%	60.000000	1117.500000	1.000000
50%	70.049958	1432.000000	2.000000
75%	78.000000	1721.000000	2.000000
max	105.000000	2626.250000	3.500000

Summary statistics after capping:

	LotFrontage	GrLivArea	FullBath
count	1459.000000	1459.000000	1459.000000
mean	68.423126	1478.000685	1.569568
std	17.164354	457.873870	0.549778
min	33.000000	407.000000	0.000000
25%	60.000000	1117.500000	1.000000
50%	70.049958	1432.000000	2.000000
75%	78.000000	1721.000000	2.000000
max	105.000000	2626.250000	3.500000

```
def cap_outliers_verbose(df, columns):
    for column in columns:
        Q1 = df[column].quantile(0.25)
        Q3 = df[column].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        print(f"{column}: Lower Bound = {lower_bound}, Upper Bound = {upper_bound}")
        df[column] = df[column].apply(lambda x: lower_bound if x < lower_bound else upper_bound if x > upper_bound else x)
    return df
```

```
# Cap outliers in the specified columns of test_data with verbose output
```

```
test_data_capped_verbose = cap_outliers_verbose(test_data.copy(),
columns_to_cap_test)
```

```
# Verify the changes
```

```
print(f"Original test_data shape: {test_data.shape}")
```

```
print(f"Capped test_data shape: {test_data_capped_verbose.shape}")
```

```
MSSubClass: Lower Bound = -55.0, Upper Bound = 145.0
LotFrontage: Lower Bound = 33.0, Upper Bound = 105.0
LotArea: Lower Bound = 1201.25, Upper Bound = 17707.25
OverallQual: Lower Bound = 2.0, Upper Bound = 10.0
OverallCond: Lower Bound = 3.5, Upper Bound = 7.5
YearBuilt: Lower Bound = 1881.0, Upper Bound = 2073.0
MasVnrArea: Lower Bound = -243.0, Upper Bound = 405.0
BsmtFinSF1: Lower Bound = -1128.0, Upper Bound = 1880.0
BsmtFinSF2: Lower Bound = 0.0, Upper Bound = 0.0
BsmtUnfSF: Lower Bound = -647.5, Upper Bound = 1664.5
TotalBsmtSF: Lower Bound = 4.0, Upper Bound = 2084.0
1stFlrSF: Lower Bound = 110.0, Upper Bound = 2146.0
2ndFlrSF: Lower Bound = -1014.0, Upper Bound = 1690.0
LowQualFinSF: Lower Bound = 0.0, Upper Bound = 0.0
GrLivArea: Lower Bound = 212.25, Upper Bound = 2626.25
BsmtFullBath: Lower Bound = -1.5, Upper Bound = 2.5
BsmtHalfBath: Lower Bound = 0.0, Upper Bound = 0.0
FullBath: Lower Bound = -0.5, Upper Bound = 3.5
BedroomAbvGr: Lower Bound = 0.5, Upper Bound = 4.5
KitchenAbvGr: Lower Bound = 1.0, Upper Bound = 1.0
TotRmsAbvGrd: Lower Bound = 2.0, Upper Bound = 10.0
Fireplaces: Lower Bound = -1.5, Upper Bound = 2.5
GarageYrBlt: Lower Bound = 1899.75, Upper Bound = 2061.75
GarageCars: Lower Bound = -0.5, Upper Bound = 3.5
GarageArea: Lower Bound = -69.0, Upper Bound = 963.0
WoodDeckSF: Lower Bound = -252.0, Upper Bound = 420.0
OpenPorchSF: Lower Bound = -108.0, Upper Bound = 180.0
EnclosedPorch: Lower Bound = 0.0, Upper Bound = 0.0
3SsnPorch: Lower Bound = 0.0, Upper Bound = 0.0
ScreenPorch: Lower Bound = 0.0, Upper Bound = 0.0
PoolArea: Lower Bound = 0.0, Upper Bound = 0.0
MiscVal: Lower Bound = 0.0, Upper Bound = 0.0
Original test_data shape: (1459, 80)
Capped test_data shape: (1459, 80)
```

```
def check_outliers(df, columns):
    outliers = {}
    for column in columns:
        Q1 = df[column].quantile(0.25)
        Q3 = df[column].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
```

```

        outlier_indices = df[(df[column] < lower_bound) | (df[column]
> upper_bound)].index
        outliers[column] = len(outlier_indices)
    return outliers

```

```

# Check for outliers in the specified columns of test_data
outliers_before_capping = check_outliers(test_data,
columns_to_cap_test)

```

```

# Print outliers in the desired format
print(f"{'Column Name':<25} {'Outliers Found':<15}")
for column, count in outliers_before_capping.items():
    print(f"{column:<25} {count:<15}")

```

Column Name	Outliers Found
MSSubClass	0
LotFrontage	0
LotArea	0
OverallQual	0
OverallCond	0
YearBuilt	0
MasVnrArea	0
BsmtFinSF1	0
BsmtFinSF2	0
BsmtUnfSF	0
TotalBsmtSF	0
1stFlrSF	0
2ndFlrSF	0
LowQualFinSF	0
GrLivArea	0
BsmtFullBath	0
BsmtHalfBath	0
FullBath	0
BedroomAbvGr	0
KitchenAbvGr	0
TotRmsAbvGrd	0
Fireplaces	0
GarageYrBlt	0
GarageCars	0
GarageArea	0
WoodDeckSF	0
OpenPorchSF	0
EnclosedPorch	0
3SsnPorch	0
ScreenPorch	0
PoolArea	0
MiscVal	0

```

def cap_outliers_verbose(df, columns):
    for column in columns:

```

```

    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    print(f"{column}: Lower Bound = {lower_bound}, Upper Bound = {upper_bound}")
    df[column] = df[column].apply(lambda x: lower_bound if x < lower_bound else upper_bound if x > upper_bound else x)
    return df

```

Cap outliers in the specified columns of test_data with verbose output

```
train_data_capped_verbose = cap_outliers_verbose(train_data.copy(), columns_to_cap_test)
```

Verify the changes

```
print(f"Original test_data shape: {train_data.shape}")
print(f"Capped test_data shape: {train_data_capped_verbose.shape}")

```

```

MSSubClass: Lower Bound = -55.0, Upper Bound = 145.0
LotFrontage: Lower Bound = 31.5, Upper Bound = 107.5
LotArea: Lower Bound = 1481.5, Upper Bound = 17673.5
OverallQual: Lower Bound = 2.0, Upper Bound = 10.0
OverallCond: Lower Bound = 3.5, Upper Bound = 7.5
YearBuilt: Lower Bound = 1885.0, Upper Bound = 2069.0
MasVnrArea: Lower Bound = -246.375, Upper Bound = 410.625
BsmtFinSF1: Lower Bound = -1068.375, Upper Bound = 1780.625
BsmtFinSF2: Lower Bound = 0.0, Upper Bound = 0.0
BsmtUnfSF: Lower Bound = -654.5, Upper Bound = 1685.5
TotalBsmtSF: Lower Bound = 42.0, Upper Bound = 2052.0
1stFlrSF: Lower Bound = 118.125, Upper Bound = 2155.125
2ndFlrSF: Lower Bound = -1092.0, Upper Bound = 1820.0
LowQualFinSF: Lower Bound = 0.0, Upper Bound = 0.0
GrLivArea: Lower Bound = 158.625, Upper Bound = 2747.625
BsmtFullBath: Lower Bound = -1.5, Upper Bound = 2.5
BsmtHalfBath: Lower Bound = 0.0, Upper Bound = 0.0
FullBath: Lower Bound = -0.5, Upper Bound = 3.5
BedroomAbvGr: Lower Bound = 0.5, Upper Bound = 4.5
KitchenAbvGr: Lower Bound = 1.0, Upper Bound = 1.0
TotRmsAbvGrd: Lower Bound = 2.0, Upper Bound = 10.0
Fireplaces: Lower Bound = -1.5, Upper Bound = 2.5
GarageYrBlt: Lower Bound = 1903.5, Upper Bound = 2059.5
GarageCars: Lower Bound = -0.5, Upper Bound = 3.5
GarageArea: Lower Bound = -27.75, Upper Bound = 938.25
WoodDeckSF: Lower Bound = -252.0, Upper Bound = 420.0
OpenPorchSF: Lower Bound = -102.0, Upper Bound = 170.0
EnclosedPorch: Lower Bound = 0.0, Upper Bound = 0.0
3SsnPorch: Lower Bound = 0.0, Upper Bound = 0.0
ScreenPorch: Lower Bound = 0.0, Upper Bound = 0.0

```

PoolArea: Lower Bound = 0.0, Upper Bound = 0.0

MiscVal: Lower Bound = 0.0, Upper Bound = 0.0

Original test_data shape: (1460, 81)

Capped test_data shape: (1460, 81)

```
def check_outliers(df, columns):
    outliers = {}
    for column in columns:
        Q1 = df[column].quantile(0.25)
        Q3 = df[column].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        outlier_indices = df[(df[column] < lower_bound) | (df[column]
> upper_bound)].index
        outliers[column] = len(outlier_indices)
    return outliers
```

```
# Check for outliers in the specified columns of test_data
outliers_before_capping = check_outliers(train_data,
columns_to_cap_test)
```

```
# Print outliers in the desired format
print(f"{'Column Name':<25} {'Outliers Found':<15}")
for column, count in outliers_before_capping.items():
    print(f"{column:<25} {count:<15}")
```

Column Name	Outliers Found
MSSubClass	103
LotFrontage	106
LotArea	69
OverallQual	2
OverallCond	125
YearBuilt	7
MasVnrArea	98
BsmtFinSF1	7
BsmtFinSF2	167
BsmtUnfSF	29
TotalBsmtSF	61
1stFlrSF	20
2ndFlrSF	2
LowQualFinSF	26
GrLivArea	31
BsmtFullBath	1
BsmtHalfBath	82
FullBath	0
BedroomAbvGr	35
KitchenAbvGr	68
TotRmsAbvGrd	30
Fireplaces	5

GarageYrBlt	1
GarageCars	5
GarageArea	21
WoodDeckSF	32
OpenPorchSF	77
EnclosedPorch	208
3SsnPorch	24
ScreenPorch	116
PoolArea	7
MiscVal	52

Correlation Analysis

to accurately understand the relationships between features, leading to better feature selection and model performance.

```
import pandas as pd

# Add SalePrice to the capped DataFrame for correlation analysis
train_data_capped_verbose['SalePrice'] = train_data['SalePrice']

# Calculate the correlation matrix
correlation_matrix = train_data_capped_verbose.corr()

# Extract the correlation with SalePrice
correlation_with_saleprice = correlation_matrix['SalePrice']

# Display the correlation values
print(correlation_with_saleprice)
```

Id	-0.021917
MSSubClass	-0.063602
MSZoning	-0.166872
LotFrontage	0.371558
LotArea	0.432216
Street	0.041036
Alley	0.037646
LotShape	-0.255580
LandContour	0.015453
Utilities	-0.014314
LotConfig	-0.067396
LandSlope	0.051152
Neighborhood	0.009118
Condition1	0.091155
Condition2	0.007513
BldgType	-0.084931
HouseStyle	0.206210
OverallQual	0.791965

OverallCond	-0.106261
YearBuilt	0.524172
YearRemodAdd	0.507101
RoofStyle	0.222405
RoofMatl	0.132383
Exterior1st	-0.163240
Exterior2nd	-0.103815
MasVnrType	0.113287
MasVnrArea	0.452127
ExterQual	-0.636884
ExterCond	0.117303
Foundation	0.382479
BsmtQual	-0.581349
BsmtCond	0.065844
BsmtExposure	-0.276932
BsmtFinType1	-0.072068
BsmtFinSF1	0.400330
BsmtFinType2	0.037640
BsmtFinSF2	NaN
BsmtUnfSF	0.203278
TotalBsmtSF	0.636999
Heating	-0.098812
HeatingQC	-0.400178
CentralAir	0.251328
Electrical	0.234760
1stFlrSF	0.620743
2ndFlrSF	0.316547
LowQualFinSF	NaN
GrLivArea	0.708153
BsmtFullBath	0.227813
BsmtHalfBath	NaN
FullBath	0.560664
HalfBath	0.284108
BedroomAbvGr	0.185740
KitchenAbvGr	NaN
KitchenQual	-0.589189
TotRmsAbvGrd	0.536067
Functional	0.115328
Fireplaces	0.468700
FireplaceQu	-0.079038
GarageType	-0.321603
GarageYrBlt	0.470269
GarageFinish	-0.485097
GarageCars	0.644002
GarageArea	0.630138
GarageQual	0.107677
GarageCond	0.135823
PavedDrive	0.231357
WoodDeckSF	0.330378


```
OpenPorchSF      0.369024
EnclosedPorch    NaN
3SsnPorch        NaN
ScreenPorch      NaN
PoolArea         NaN
PoolQC           -0.054580
Fence            -0.019741
MiscFeature      0.012167
MiscVal          NaN
MoSold           0.046432
YrSold           -0.028923
SaleType         -0.054911
SaleCondition    0.213092
SalePrice        1.000000
Name: SalePrice, dtype: float64
```

```
train_data_capped_verbose.isnull().sum()
```

```
Id                0
MSSubClass        0
MSZoning          0
LotFrontage       0
LotArea           0
Street            0
Alley             0
LotShape          0
LandContour       0
Utilities         0
LotConfig         0
LandSlope         0
Neighborhood      0
Condition1        0
Condition2        0
BldgType          0
HouseStyle        0
OverallQual       0
OverallCond       0
YearBuilt         0
YearRemodAdd      0
RoofStyle         0
RoofMatl          0
Exterior1st       0
Exterior2nd       0
MasVnrType        0
MasVnrArea        0
ExterQual         0
ExterCond         0
Foundation        0
```

BsmtQual	0
BsmtCond	0
BsmtExposure	0
BsmtFinType1	0
BsmtFinSF1	0
BsmtFinType2	0
BsmtFinSF2	0
BsmtUnfSF	0
TotalBsmtSF	0
Heating	0
HeatingQC	0
CentralAir	0
Electrical	0
1stFlrSF	0
2ndFlrSF	0
LowQualFinSF	0
GrLivArea	0
BsmtFullBath	0
BsmtHalfBath	0
FullBath	0
HalfBath	0
BedroomAbvGr	0
KitchenAbvGr	0
KitchenQual	0
TotRmsAbvGrd	0
Functional	0
Fireplaces	0
FireplaceQu	0
GarageType	0
GarageYrBlt	0
GarageFinish	0
GarageCars	0
GarageArea	0
GarageQual	0
GarageCond	0
PavedDrive	0
WoodDeckSF	0
OpenPorchSF	0
EnclosedPorch	0
3SsnPorch	0
ScreenPorch	0
PoolArea	0
PoolQC	0
Fence	0
MiscFeature	0
MiscVal	0
MoSold	0
YrSold	0
SaleType	0

```
SaleCondition    0
SalePrice        0
dtype: int64
```

Dropping Columns with constant Values (A column with constant values has the same value for every row in the dataset. Example: If a column X has the value '20' for all rows, it is a constant column.

from the correlation results, found out the data contains columns with constant values which should be removed because

No Variability: Since the values do not change, the column does not provide any useful information for analysis.

Correlation: The correlation of a constant column with any other variable is undefined or zero because there is no variability to compare.

Model Performance: Including constant columns in a model does not add any predictive power and can be safely removed.

```
# Identify columns with constant values in train_data_capped_verbose
constant_columns_train = [col for col in
train_data_capped_verbose.columns if
train_data_capped_verbose[col].nunique() == 1]
```

```
# Display the constant columns
print("Columns with constant values in train_data_capped_verbose:",
constant_columns_train)
```

```
Columns with constant values in train_data_capped_verbose:
['BsmtFinSF2', 'LowQualFinSF', 'BsmtHalfBath', 'KitchenAbvGr',
'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal']
```

```
# Drop columns with constant values from train_data_capped_verbose
train_data_capped_verbose.drop(columns=constant_columns_train,
inplace=True)
```

```
# Verify the changes
print("Remaining columns in train_data_capped_verbose:",
train_data_capped_verbose.columns)
```

```
Remaining columns in train_data_capped_verbose: Index(['Id',
'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',
'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',
'LandSlope', 'Neighborhood', 'Condition1', 'Condition2',
'BldgType',
'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt',
'YearRemodAdd',
'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd',
'MasVnrType',
```

```

        'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation',
        'BsmtQual',
        'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1',
        'BsmtFinType2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating',
        'HeatingQC',
        'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF',
        'GrLivArea',
        'BsmtFullBath', 'FullBath', 'HalfBath', 'BedroomAbvGr',
        'KitchenQual',
        'TotRmsAbvGrd', 'Functional', 'Fireplaces', 'FireplaceQu',
        'GarageType',
        'GarageYrBlt', 'GarageFinish', 'GarageCars', 'GarageArea',
        'GarageQual',
        'GarageCond', 'PavedDrive', 'WoodDeckSF', 'OpenPorchSF',
        'PoolQC',
        'Fence', 'MiscFeature', 'MoSold', 'YrSold', 'SaleType',
        'SaleCondition',
        'SalePrice'],
        dtype='object')

```

```
train_data_capped_verbose.shape
```

```
(1460, 72)
```

```
# Identify columns with constant values
```

```
constant_columns = [col for col in test_data_capped_verbose.columns if
test_data_capped_verbose[col].nunique() == 1]
```

```
# Display the constant columns
```

```
print("Columns with constant values:", constant_columns)
```

```
Columns with constant values: ['Utilities', 'BsmtFinSF2',
'LowQualFinSF', 'BsmtHalfBath', 'KitchenAbvGr', 'EnclosedPorch',
'3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal']
```

```
# Drop columns with constant values
```

```
test_data_capped_verbose.drop(columns=constant_columns, inplace=True)
```

```
# Verify the changes
```

```
print("Remaining columns:", test_data_capped_verbose.columns)
```

```
Remaining columns: Index(['Id', 'MSSubClass', 'MSZoning',
'LotFrontage', 'LotArea', 'Street',
'Alley', 'LotShape', 'LandContour', 'LotConfig', 'LandSlope',
'Neighborhood', 'Condition1', 'Condition2', 'BldgType',
'HouseStyle',
'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd',
'RoofStyle',
'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType',
'MasVnrArea',
'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual', 'BsmtCond',
```

```

        'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1', 'BsmtFinType2',
        'BsmtUnfSF', 'TotalBsmtSF', 'Heating', 'HeatingQC',
        'CentralAir',
        'Electrical', '1stFlrSF', '2ndFlrSF', 'GrLivArea',
        'BsmtFullBath',
        'FullBath', 'HalfBath', 'BedroomAbvGr', 'KitchenQual',
        'TotRmsAbvGrd',
        'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType',
        'GarageYrBlt',
        'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual',
        'GarageCond',
        'PavedDrive', 'WoodDeckSF', 'OpenPorchSF', 'PoolQC', 'Fence',
        'MiscFeature', 'MoSold', 'YrSold', 'SaleType',
        'SaleCondition'],
        dtype='object')

test_data_capped_verbose.shape

(1459, 70)

```

Dropping columns with low correlation to the Target Variable (SalePrice)

Correlation Strength:: Strong Correlation: Typically, correlations above 0.5 (positive or negative)

```

import pandas as pd

# Calculate the correlation matrix for the train dataset
corr_matrix = train_data_capped_verbose.corr()

# Set the correlation threshold
threshold = 0.5

# Get features with a correlation above or below the threshold for the
train dataset
correlated_features = corr_matrix.index[(corr_matrix["SalePrice"] >=
threshold) | (corr_matrix["SalePrice"] <= -threshold)]

# Filter the train dataset
train_data_filtered = train_data_capped_verbose[correlated_features]

# Filter the test dataset using the same correlated features
(excluding 'SalePrice' if it exists)
correlated_features_test = correlated_features.drop('SalePrice',
errors='ignore')
test_data_filtered =
test_data_capped_verbose[correlated_features_test]

```

```
# Display the first few rows of the filtered train and test datasets
print(train_data_filtered.head())
print(test_data_filtered.head())
```

	OverallQual	YearBuilt	YearRemodAdd	ExterQual	BsmtQual
0	7.0	2003.0	2003.0	2	2
1	6.0	1976.0	1976.0	3	2
2	7.0	2001.0	2002.0	2	2
3	7.0	1915.0	1970.0	3	3
4	8.0	2000.0	2000.0	2	2

	1stFlrSF	GrLivArea	FullBath	KitchenQual	TotRmsAbvGrd
0	856.0	1710.0	2.0	2	8.0
1	1262.0	1262.0	2.0	3	6.0
2	920.0	1786.0	2.0	2	6.0
3	961.0	1717.0	1.0	2	7.0
4	1145.0	2198.0	2.0	2	9.0

	GarageArea	SalePrice
0	548.0	208500
1	460.0	181500
2	608.0	223500
3	642.0	140000
4	836.0	250000

	OverallQual	YearBuilt	YearRemodAdd	ExterQual	BsmtQual
0	5.0	1961.0	1961.0	7	7
1	6.0	1958.0	1958.0	7	7
2	5.0	1997.0	1998.0	7	6
3	6.0	1998.0	1998.0	7	7
4	8.0	1992.0	1992.0	6	6

	1stFlrSF	GrLivArea	FullBath	KitchenQual	TotRmsAbvGrd
--	----------	-----------	----------	-------------	--------------

```

GarageCars \
0      896.0      896.0      1.0      7      5.0
1.0
1      1329.0     1329.0      1.0      6      6.0
1.0
2      928.0      1629.0      2.0      7      6.0
2.0
3      926.0      1604.0      2.0      6      7.0
2.0
4      1280.0     1280.0      2.0      6      5.0
2.0

```

```

GarageArea
0      730.0
1      312.0
2      482.0
3      470.0
4      506.0

```

```
train_data_filtered.shape
```

```
(1460, 14)
```

```
test_data_filtered.shape
```

```
(1459, 13)
```

```
train_data_filtered.head()
```

```

OverallQual  YearBuilt  YearRemodAdd  ExterQual  BsmtQual
TotalBsmtSF \
0      7.0      2003.0      2003.0      2      2
856.0
1      6.0      1976.0      1976.0      3      2
1262.0
2      7.0      2001.0      2002.0      2      2
920.0
3      7.0      1915.0      1970.0      3      3
756.0
4      8.0      2000.0      2000.0      2      2
1145.0

```

```

1stFlrSF  GrLivArea  FullBath  KitchenQual  TotRmsAbvGrd
GarageCars \
0      856.0      1710.0      2.0      2      8.0
2.0
1      1262.0     1262.0      2.0      3      6.0
2.0
2      920.0      1786.0      2.0      2      6.0
2.0
3      961.0      1717.0      1.0      2      7.0

```

3.0					
4	1145.0	2198.0	2.0	2	9.0
3.0					

	GarageArea	SalePrice
0	548.0	208500
1	460.0	181500
2	608.0	223500
3	642.0	140000
4	836.0	250000

Scaling the data and Splitting the data

1. **Equal Contribution of Features** Scaling ensures that all features contribute equally to the model. Without scaling, features with larger ranges can dominate the model's performance, leading to biased results¹.
2. **Improved Model Performance** Many machine learning algorithms, including Linear Regression, perform better when the data is scaled. This is because these algorithms often rely on distance calculations, and unscaled data can distort these distances¹.
3. **Faster Convergence** Scaling can speed up the convergence of gradient descent, which is used in many optimization algorithms. This results in faster training times and more efficient model building¹.
4. **Consistency in Data Scaling** helps maintain consistency in the data, making it easier to compare and interpret the results. This is particularly important when dealing with features that have different units or scales¹.

```
from sklearn.preprocessing import StandardScaler
import pandas as pd

# Initialize the scaler
scaler = StandardScaler()

# Separate the features (X) and target (y)
X_train = train_data_filtered.drop(columns=['SalePrice'])
y_train = train_data_filtered['SalePrice']

# Standardize the numeric features in the training set
X_train_scaled = pd.DataFrame(scaler.fit_transform(X_train),
                              columns=X_train.columns)

# Standardize the numeric features in the test set
X_test_scaled = pd.DataFrame(scaler.transform(test_data_filtered),
                              columns=X_train.columns) # Use X_train's columns for consistency

# Check the first few rows of the standardized training data
print(X_train_scaled.head())
```


	OverallQual	YearBuilt	YearRemodAdd	ExterQual	BsmtQual	
TotalBsmtSF \						
0	0.652644	1.053246	0.878668	-0.777976	-0.290552	-
0.488321						
1	-0.073068	0.156179	-0.429577	0.663451	-0.290552	
0.532289						
2	0.652644	0.986797	0.830215	-0.777976	-0.290552	-
0.327437						
3	0.652644	-1.870528	-0.720298	0.663451	0.852861	-
0.739702						
4	1.378355	0.953572	0.733308	-0.777976	-0.290552	
0.238172						

	1stFlrSF	GrLivArea	FullBath	KitchenQual	TotRmsAbvGrd
GarageCars \					
0	-0.830489	0.428636	0.789741	-0.409369	0.981148
0.315946					
1	0.289638	-0.502349	0.789741	0.795629	-0.316385
0.315946					
2	-0.653917	0.586571	0.789741	-0.409369	-0.316385
0.315946					
3	-0.540801	0.443182	-1.026041	-0.409369	0.332382
1.662750					
4	-0.033157	1.442744	0.789741	-0.409369	1.629914
1.662750					

	GarageArea
0	0.373509
1	-0.051541
2	0.663315
3	0.827539
4	1.764579

Importing the simple Linear Regression model for the FIT, training and testing .

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Initialize the Linear Regression model
model = LinearRegression()

# Train the model on standardized data
model.fit(X_train_scaled, y_train)
```

```

# Predict the test data
y_pred = model.predict(X_test_scaled)

from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error
import numpy as np

# Assuming your features and target variable are correctly assigned
X_train = X_train_scaled  ##.drop(columns=['SalePrice'])
y_train = train_data['SalePrice']

# Initialize and train the Linear Regression model
model = LinearRegression()
model.fit(X_train, y_train)

LinearRegression()

```

Preparing the Test data set for Prediction

```

X_test = X_test_scaled

y_test_pred = model.predict(X_test)

print("Predicted house prices for the test data:")
print(y_test_pred)

Predicted house prices for the test data:
[ 3581.49894512  57894.581041   57157.77183062 ...  41187.92002634
 4818.91918207 114638.14467042]

```

Model performance Evaluation Using Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), R-squared and Mean Squared Error (MSE) metrics

```

from sklearn.metrics import mean_absolute_error, mean_squared_error

# Predictions on the training data
y_train_pred = model.predict(X_train)

```

```

# Calculate MAE and RMSE
mae = mean_absolute_error(y_train, y_train_pred)
rmse = mean_squared_error(y_train, y_train_pred, squared=False)

# Print the results
print(f'Training MAE: {mae}')
print(f'Training RMSE: {rmse}')

# Additional metrics
r_squared = model.score(X_train, y_train)
mse = mean_squared_error(y_train, y_train_pred)

print(f'R-squared on training data: {r_squared}')
print(f'Mean Squared Error on training data: {mse}')

Training MAE: 22423.206122427106
Training RMSE: 34668.73596670487
R-squared on training data: 0.8094242042104971
Mean Squared Error on training data: 1201921253.5290961

```

Mean and Median of the training dataset

```

mean_price = train_data['SalePrice'].mean()
median_price = train_data['SalePrice'].median()
print(f"Mean SalePrice: {mean_price}")
print(f"Median SalePrice: {median_price}")

Mean SalePrice: 180921.19589041095
Median SalePrice: 163000.0

# Make predictions on the training data
train_predictions = model.predict(X_train)

# Print first 5 predictions alongside actual prices
for i in range(30):
    print(f'Actual SalePrice: {y_train.iloc[i]}, Predicted SalePrice: {train_predictions[i]}')

Actual SalePrice: 208500, Predicted SalePrice: 212060.88445107004
Actual SalePrice: 181500, Predicted SalePrice: 163409.79247865872
Actual SalePrice: 223500, Predicted SalePrice: 220740.0983045473
Actual SalePrice: 140000, Predicted SalePrice: 189051.0719786906
Actual SalePrice: 250000, Predicted SalePrice: 276036.7465114494
Actual SalePrice: 143000, Predicted SalePrice: 149592.65107443393
Actual SalePrice: 307000, Predicted SalePrice: 282254.5452889392
Actual SalePrice: 200000, Predicted SalePrice: 216356.6533658756
Actual SalePrice: 129900, Predicted SalePrice: 172366.86374207088
Actual SalePrice: 118000, Predicted SalePrice: 106490.10264236729
Actual SalePrice: 129500, Predicted SalePrice: 116711.88445916242

```

```
Actual SalePrice: 345000, Predicted SalePrice: 342315.6425546423
Actual SalePrice: 144000, Predicted SalePrice: 102546.95500961428
Actual SalePrice: 279500, Predicted SalePrice: 240452.01629296137
Actual SalePrice: 157000, Predicted SalePrice: 148967.28469024325
Actual SalePrice: 132000, Predicted SalePrice: 133607.35428353102
Actual SalePrice: 149000, Predicted SalePrice: 135015.75113638488
Actual SalePrice: 90000, Predicted SalePrice: 88617.97272179213
Actual SalePrice: 159000, Predicted SalePrice: 154779.91973603086
Actual SalePrice: 139000, Predicted SalePrice: 132587.38890276017
Actual SalePrice: 325300, Predicted SalePrice: 300014.33701249905
Actual SalePrice: 139400, Predicted SalePrice: 134226.27870761263
Actual SalePrice: 230000, Predicted SalePrice: 269553.77691814094
Actual SalePrice: 129900, Predicted SalePrice: 139821.63973336632
Actual SalePrice: 154000, Predicted SalePrice: 130924.46034890861
Actual SalePrice: 256300, Predicted SalePrice: 264891.46087731625
Actual SalePrice: 134800, Predicted SalePrice: 124779.86412297426
Actual SalePrice: 306000, Predicted SalePrice: 292243.4579140794
Actual SalePrice: 207500, Predicted SalePrice: 169933.99781867318
Actual SalePrice: 68500, Predicted SalePrice: 59136.46397952481
```

Cross- Validation

used to assess the performance of a machine learning model by dividing the dataset into multiple subsets, or "folds," and then training and testing the model multiple times on different combinations of these folds. It helps evaluate how well the model generalizes to unseen data and prevents overfitting.

```
from sklearn.model_selection import cross_val_score
from sklearn.metrics import make_scorer, mean_absolute_error,
mean_squared_error
import numpy as np

# Define scorers
mae_scorer = make_scorer(mean_absolute_error)
rmse_scorer = make_scorer(mean_squared_error, squared=False)

# Perform cross-validation
mae_scores = cross_val_score(model, X_train, y_train,
                              scoring=mae_scorer, cv=5)
rmse_scores = cross_val_score(model, X_train, y_train,
                              scoring=rmse_scorer, cv=5)

print("Cross-validated MAE: ", np.mean(mae_scores))
print("Cross-validated RMSE: ", np.mean(rmse_scores))

Cross-validated MAE: 22807.962869899988
Cross-validated RMSE: 34994.64321135447
```

Model Performance Summary:

The Linear Regression model was evaluated on both the training set and through cross-validation, yielding the following results:

1. Training MAE (Mean Absolute Error): 22,423.21
2. Training RMSE (Root Mean Squared Error): 34,668.74
3. R-squared on training data: 0.8094, indicating that the model explains approximately 81% of the variance in the target variable (SalePrice).
4. Cross-validated MAE: 22,807.96
5. Cross-validated RMSE: 34,994.64

These performance metrics suggest that the model generalizes well across different subsets of the data, as indicated by the similarity between the training and cross-validated results. The MAE and RMSE values reveal that the model's predictions are, on average, off by about 22,000 in SalePrice, and the magnitude of prediction errors is around 34,000.

Insights:

1. The model explains a significant portion of the variance in SalePrice, with an R-squared score of approximately 81%. However, the error metrics (MAE and RMSE) indicate that the model's predictions are still far from perfect.
2. With a mean SalePrice of 180,921 and a median SalePrice of 163,000, the MAE and RMSE values of 22,000+ and 34,000+ represent a substantial error relative to these typical SalePrice values.
3. The closeness of the training and cross-validated error metrics suggests that the model is not overfitting and generalizes well, but there is still room for improvement to reduce the prediction errors.

Recommendations for Improvement:

1. Feature Engineering: Introduce interaction terms or non-linear features to better capture the relationships in the data.
2. Advanced Models: Consider using more sophisticated models like Random Forests, Gradient Boosting Machines (GBM), or XGBoost, which may capture complex patterns better than linear models.
3. Hyperparameter Tuning: To explore more advanced models, performing hyperparameter tuning could help optimize the model's performance and reduce prediction errors.

4. Overall, the model shows promise with consistent results across training and cross-validation, but additional improvements could further enhance prediction

